

Pre-trained language models (LMs) with billions of parameters are the go-to models for many NLP tasks. The sheer scale of their training corpora and parameter counts appears to endow them with the ability to compete with, and frequently outperform, task-specific models. LMs, however, can often have unexpected failure modes and are notoriously difficult to control. The responsible refinement and improvement of these models depends upon the creation of the necessary tools and frameworks for understanding them. While an undergraduate at Harvard (advised by Stuart Shieber and Yonatan Belinkov) and subsequently as a pre-doctoral researcher at AI2 (advised by Peter Clark and Ashish Sabharwal), I have adopted an approach of **drawing on linguistics** such as syntax and formal language theory to develop practical and theoretical frameworks towards **understanding, evaluating, and mitigating risks** from language models as **general-purpose NLP systems**.

**Understanding and improving LMs through linguistics** Modern NLP has become largely divorced from our understanding of linguistics and cognition in humans. Reconciling these two can lend clarity to our understanding of how LMs work. As my first foray into this effort, I set out to discover how modern LMs handle syntactic agreement. In my first-author ACL 2021 paper [1] we intervene on individual neurons in large transformers to observe their causal effect on syntactic agreement. We find, among other things, that transformers learn two distinct mechanisms for number agreement, and that these mechanisms are distributed in the activations of the network, rather than concentrated in any single model component, as was found with gender bias [2]. These findings constitute a step forward in being able to understand and interpret how LMs mimic syntactic behavior, shedding light into the black box model. Subsequent research has built on our work in order to interpret other linguistic phenomena such as distributivity [3]. In future research, I hope to explore the potential of linguistically informed approaches to evaluate and improve the capabilities of LMs. As an example, I am interested in leveraging the language acquisition literature to evaluate and improve LMs’ abilities to acquire new words by inferring their meaning from context. Improving this ability might be achieved via a self-supervised pre-training objective that randomly replaces lexical items with new unseen words, forcing the model to learn how to generalize effectively. Success here would result in LMs that are more robust to linguistic distribution shifts and adapt to evolving language.

**Understanding LMs through formal languages** I am excited by projects that borrow from formal language theory to increase our understanding of LMs. This approach makes it possible to answer questions that might otherwise be very difficult to approach using natural language. For instance, I used regular languages to **measure the capabilities of transformers as instruction followers** in RegSet, my first-author EMNLP 2022 paper [4]. Large, pre-trained LMs can solve some NLP tasks by conditioning their generations on natural language instructions for the task [5, 6]. On the other hand, research has shown that neural models consistently struggle with compositional generalization [7]. This bodes poorly for the instruction following regime where the space of task descriptions is both intractably large and highly compositional. Moreover, the complexity of natural language makes it difficult to predict what types of instructions may be challenging for transformers. To solve this predicament, we propose a controllable proxy for studying instruction learning by studying LMs’ ability to follow instructions in the form of regular expressions. We test the effects of attributes of regular languages, such as starfreeness, on their difficulty as instructions. Our experiments lead us to a number of intriguing hypotheses about what makes instruction learning hard, including evidence that even large transformers struggle with modular counting (e.g., determining whether something is even or odd). By taking advantage of the well studied attributes of formal languages, we achieve fine-grained control over our data, leading to findings that would have been extremely difficult and expensive to obtain on natural data. This approach can be applied more broadly to develop benchmarks that isolate and measure progress towards specific abilities in transformers that we hope to see in natural language settings.

Formal language theory can also be used to derive theoretical results that help us understand modern

neural architectures. I am currently developing a framework for understanding what transformers can learn from instructions. In particular, I hope to show **which formal language families are provably learnable by transformers via instructions**. This builds on prior work [8], and provides both a principled way to study transformers and bounds on what we can expect them to learn. Another useful direction could be to study **the theoretical implications of sub-task decomposition**. I have previously worked on empirical studies in this area: in a preprint currently under submission to ICLR 2023 [9], I implement a modular method for recursively prompting large LMs which vastly improves generalization to longer sequence lengths for certain types of tasks when compared to other step-by-step reasoning styles. Recent work [10] proves that subtask decomposition via step-by-step reasoning enables learning difficult sequence-to-sequence tasks. Extending this work by characterizing the additional computational power these techniques afford transformers would be both intellectually interesting and relevant towards understanding why step-by-step reasoning has emerged as such an effective strategy for prompting LMs.

**Evaluating and mitigating risks from general-purpose LMs** Proper evaluation is critical to advancing general-purpose NLP systems, and current methods are often insufficient for this purpose. For instance, existing datasets for are too narrow in scope to holistically evaluate general-purpose math reasoning skills in LMs. To address this, I led a team of 11 researchers in compiling a comprehensive and diverse natural language **math reasoning benchmark** called L<sub>1</sub>LA [11] (first-author paper, EMNLP 2022). We curate over 140K math problems and provide valuable annotations for reasoning via program synthesis. Our experiments show that multitask learning and augmenting the model with a Python interpreter massively improves LM performance while also providing explicit reasoning steps via generated programs. Our multitask model, Bhāskara, outperforms similarly-sized models when fine-tuning on new math reasoning tasks. Despite our modeling contributions, L<sub>1</sub>LA also shows that LMs in their current form are woefully deficient when it comes to math reasoning, and highlights the need for unified evaluations for aspiring general-purpose reasoning models. I plan to continue to create thoughtful and comprehensive evaluations to measure and promote research into models with greater general utility. In particular, I hope to align evaluation metrics with human judgements on open-ended tasks. From my work on L<sub>1</sub>LA, I experienced how evaluation becomes increasingly difficult as tasks become more open-ended. F1, exact match, and even more sophisticated metrics like Rouge, still leave much to be desired. I am interested in building *learned* metrics, i.e., using a model to assign scores, that can capture the subtle nuances required to evaluate open-ended tasks, and realign automatic evaluation with human judgement.

During my PhD I also hope to expand research on the risks associated with general-purpose NLP systems and how to mitigate them. For instance, I am currently engineering a decoding procedure for using frozen LMs as black boxes to generate their own task-specific prompts. I hope to apply this technique to study **adversarial prompts**: prompts that appear to elicit one behavior but cause the model to exhibit another. These prompts could be generated by simultaneously decoding for fluency for one task and accuracy on another. Exposing these types of vulnerabilities enables the research community to better develop secure and ethical general-purpose NLP systems.

**Why ELLIS** I place great value on autonomy in choosing my research directions. My best work comes from focusing my energy on research problems in NLP and computational linguistics that excite me intellectually and allow me to explore new ideas. ELLIS provides a unique opportunity to work with top researchers across multiple European institutions, thereby offering the multi-perspective experience wherein I can thrive. The unique exchange program and dual advisors further enrich this experience. I expect that ELLIS will set me up for success as in my academic career by opening avenues for mentorship and development during my PhD, and lasting relationships that I intend to maintain beyond into postdoctoral research, eventually helping me fulfill my goal of becoming a PI.

## References

- [1] Matthew Finlayson, Aaron Mueller, Sebastian Gehrmann, Stuart Shieber, Tal Linzen, and Yonatan Belinkov. Causal analysis of syntactic agreement mechanisms in neural language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1828–1843, Online, August 2021. Association for Computational Linguistics.
- [2] Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart M. Shieber. Investigating gender bias in language models using causal mediation analysis. In *NeurIPS*, 2020.
- [3] Pangbo Ban, Yifan Jiang, Tianran Liu, and Shane Steinert-Threlkeld. Testing pre-trained language models’ understanding of distributivity via causal mediation analysis. *ArXiv*, abs/2209.04761, 2022.
- [4] Matthew Finlayson, Kyle Richardson, Ashish Sabharwal, and Peter Clark. What makes instruction learning hard? an investigation and a new challenge in a synthetic environment. *EMNLP*, abs/2204.09148, 2022.
- [5] Swaroop Mishra, Daniel Khashabi, Chitta Baral, and Hannaneh Hajishirzi. Cross-task generalization via natural language crowdsourcing instructions. In *ACL*, 2022.
- [6] Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V Le. Finetuned language models are zero-shot learners. In *ICLR*, 2022.
- [7] Brenden M. Lake and Marco Baroni. Generalization without systematicity: On the compositional skills of sequence-to-sequence recurrent networks. In *ICML*, 2018.
- [8] William Merrill and Ashish Sabharwal. Log-precision transformers are constant-depth uniform threshold circuits. *ArXiv*, abs/2207.00729, 2022.
- [9] Tushar Khot, H. Trivedi, Matthew Finlayson, Yao Fu, Kyle Richardson, Peter Clark, and Ashish Sabharwal. Decomposed prompting: A modular approach for solving complex tasks. *ArXiv*, abs/2210.02406, 2022.
- [10] Noam Wies, Yoav Levine, and Amnon Shashua. Sub-task decomposition enables learning in sequence to sequence tasks. *ArXiv*, abs/2204.02892, 2022.
- [11] Swaroop Mishra, Matthew Finlayson, Pan Lu, Leonard Tang, Sean Welleck, Chitta Baral, Tanmay Rajpurohit, Oyvind Tafjord, Ashish Sabharwal, Peter Clark, and A. Kalyan. Lila: A unified benchmark for mathematical reasoning. *EMNLP*, abs/2210.17517, 2022.